

DISTRIBUTED SYSTEMS

BSc Informatics, 4th Semester



OBJECTIVES

- Describe **foundational concepts** and challenges of distributed systems
- Evaluate **software paradigms** and strategies for distributed applications
- Assess the suitability of **cloud computing** services for given distributed applications
- *Conceptualize, design and implement* **distributed software applications**
- *Integrate and deploy* a distributed application with a **cloud computing platform** (Google App Engine, Amazon EC2, Microsoft Azure)
- *Evaluate and improve* distributed solutions in regards to **security** and **performance**

CONTENTS

- Distributed processes and communication technologies (sockets, RPC/RMI)
- Principles of naming and discovery
- Fault tolerance, process resilience, consistency and replication
- Synchronization (clocks, fundamental algorithms for synchronization)
- Security in distributed systems
- Distributed Applications (peer-to-peer networking, SaaS, cloud computing, distributed scientific computing, Google App Engine, Amazon EC2, Apache Hadoop)
- Distributed file systems

COMPUTER SYSTEMS

2. SEM

Networking I

3. SEM


Networking II

Operating
Systems


4. SEM

Distributed
Systems

REFERENCES



Distributed Systems: Principles and Paradigms
3rd edition, 2017
Andrew S. Tanenbaum
Maarten Van Steen
Prentice Hall



Distributed and Cloud Computing
2011
Kai Hwang
Jack Dongarra
Geoffrey Fox

OUTLINE

1. Introduction
 - Definition of distributed system
 - Types of distributed systems
 - Distribution transparency
2. Distributed systems architecture
 - Centralized
 - Client/Server architectures
 - Decentralized
 - Peer-to-peer systems
 - Hybrid architectures

OUTLINE

3. Service Oriented Architectures (SOA)

- REST
- Web services
- Microservices
- Message-oriented Middleware

(+ Observability guest lecture 19/04)



OUTLINE

4. Cloud and Scientific Computing

- Virtualization
- Cloud computing
- Containerized applications
- Scientific computing
 - Computing grids
- Programming paradigms for parallel distributed computing
 - MapReduce
 - MPI

OUTLINE

- 5. Processes and communication
 - Processes and threads
 - Client/Server perspective
 - Threads
 - Communication
 - RPC/RMI
 - Message-oriented communication

OUTLINE

6. Naming

- Flat naming
- Structured naming
- Attributed-based naming

7. Time and synchronization

- Clock synchronization
- Logical clocks
- Mutual exclusion
- Election algorithms

OUTLINE

8. Consistency and Replication

- Data-centric consistency
- Client-centric consistency
- Replica management
- Consistency protocols

9. Resilience and Fault Tolerance

- Process resilience
- Reliable client/server communication
- Failure recovery

MODULE ASSESSMENT

- 50 points: final examination (theory)
- 50 points: exercises portfolio
 - 5 points: architecture
 - Due date: 15/03
 - 15 points: programming exercise (Web services and REST)
 - Due date: 12/04
 - 15 points: programming exercise (message queues)
 - Due date: 17/05
 - 15 points: programming exercise (fault tolerance, consistency and security)
 - Due date: 07/06
- Each exercise is based upon the previous one!

EXERCISES

- Class exercises (programming) and homework
- Deliverables (graded)
 - Pairs
 - MS Teams submission
- Companion repository: https://github.com/IMC-UAS-Krems/DS_Examples
- Python examples
- Deliverables can be in a programming language of your choice!

LECTURE OUTLINE

- Definition
- Goals of distributed systems
- Types of distributed systems
 - Examples

WHAT IS A DISTRIBUTED SYSTEM?


Quiz: Which words would you associate with a distributed system?



<https://ahaslides.com/IMCDSI>

“A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.”

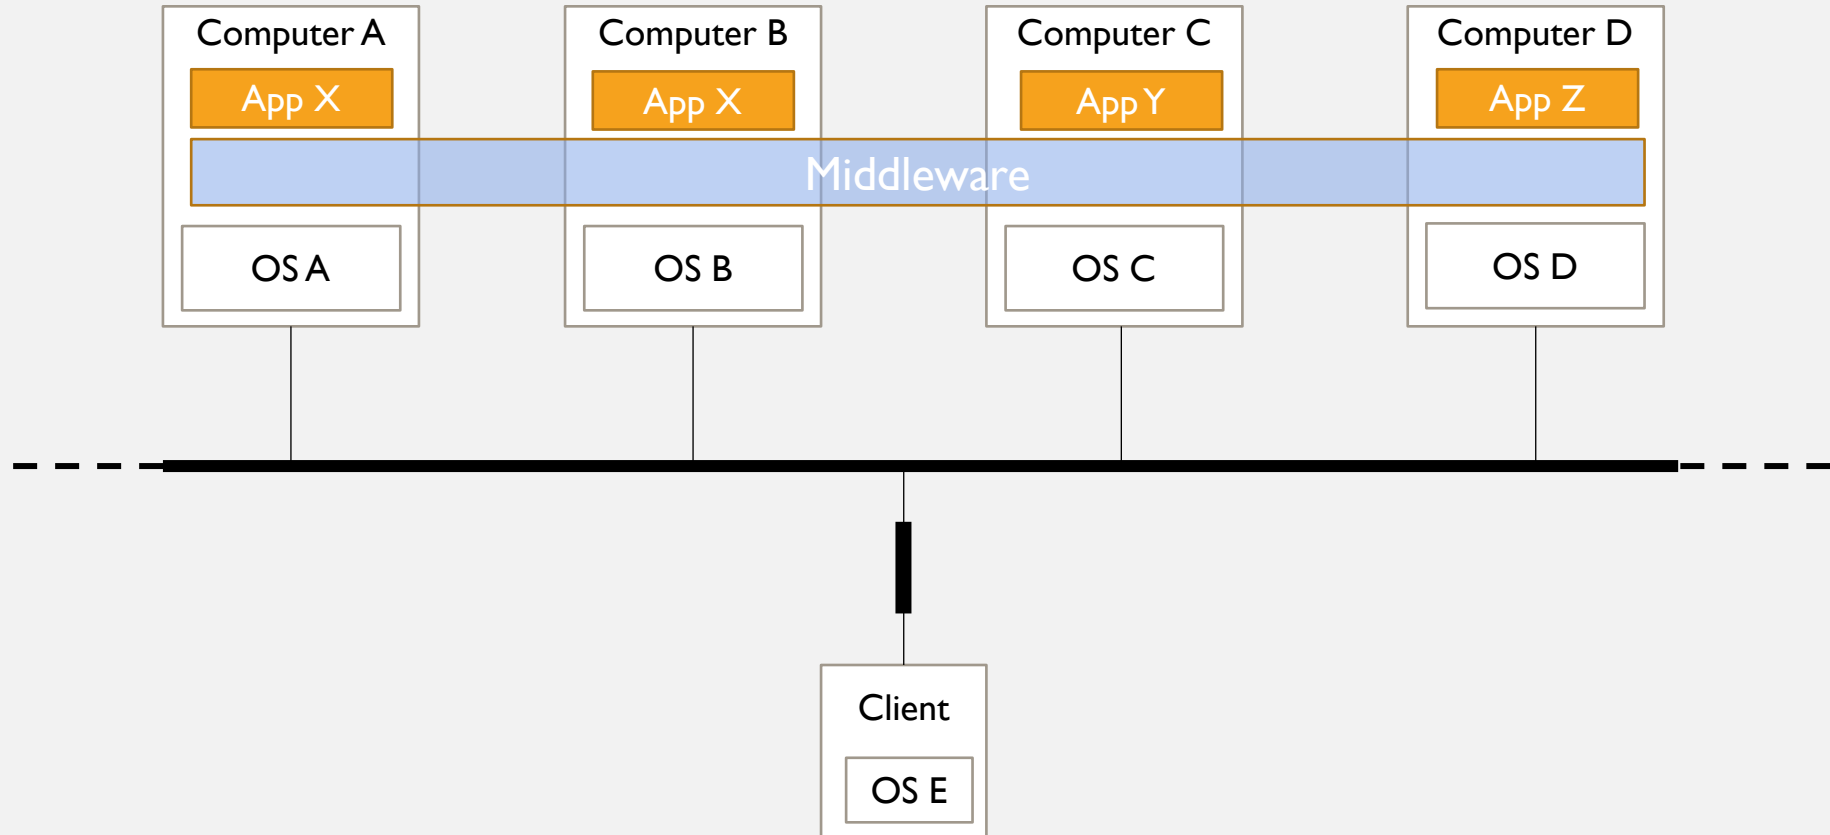
- *Leslie Lamport*



WHAT IS A DISTRIBUTED SYSTEM?

- Many definitions
- "A distributed system is a *collection of independent computers* that appears to its users as a *single coherent system*" (Tannenbaum and Van Maarten, 2007)
- Components: autonomous computers
- Coherent system → need to collaborate
 - Central aspect of developing distributed systems
- Transparency

WHAT IS A DISTRIBUTED SYSTEM?



DEFINITIONS

Middleware

- Layer between OS and application
- Makes communication possible between applications/computers
- Supplies additional services to the OS layer
- Types
 - Content-oriented: Remote Procedure Call (RPC), Remote Method Invocation (RMI)
 - Message-oriented Middleware (MOM): messaging queuing frameworks, ActiveMQ, Kafka
 - Application oriented: Web servers, Application servers, i.e. Tomcat
 - Data-oriented: caching, storage, i.e. Ignite, SQL data access
 - ...

WHY DISTRIBUTED SYSTEMS?

Quiz: Why do we need distributed systems?



WHY DISTRIBUTED SYSTEMS?

Challenges posed by distributed systems:

- Concurrency problems
 - Race conditions
 - Synchronization
- Partial failures
- How to achieve parallelism

WHY DISTRIBUTED SYSTEMS?

Goals to be met by distributed systems:

- Resource sharing
- Transparency
- Openness
- Scalability

RESOURCE SHARING

- Connect **shared resources**
 - Hardware: printers, scanners, etc.
 - Data: files, databases, web pages
 - Networks
- Main focus: **collaboration**
- Security
 - Data protection
 - Already in design phase as non-functional requirement



TRANSPARENCY

- Distributed system should appear as a single coherent system to the user
- **Transparency:** Way of abstracting details/complexity away from the user
- Different types of transparency
 - Access
 - Location
 - Migration
 - Relocation
 - Replication
 - Concurrency
 - Failure

TRANSPARENCY

Access transparency

- Addresses differences in data representation
- Hides differences between machines and operating systems
 - File naming conventions
- Hides differences between local and remote access (i.e. printers)
 - Same syntax
 - Different semantics (accessing a remote resource is more complex)

TRANSPARENCY

Location transparency

- **Hides** differences in **location** of a resource
- Use of logical names for resources
 - i.e. Use of URLs (mycompany.com/about_us)
 - No hint of where the resource is actually stored

Migration transparency

- Resource can be **moved without affecting** how they are **accessed**
 - i.e. Change location of web page while keeping URL the same

TRANSPARENCY

Relocation transparency

- **Resource** can be **moved** while in use **without** the user **noticing**
- Changing location of a file while in access
- Changing wireless access point while in motion

Replication transparency

- Hides that there are **several copies** of a given resource
- Replication transparency → Location transparency

TRANSPARENCY

Concurrency transparency

- **Hides** that access to a resource is being **shared by competing users**
- Concurrent printing of a document
 - Users A and B print a document. B sends the request shortly after A. If there is no transparency, the printer would print both documents interleaved
- Resource should remain in a consistent state
- Locks/Synchronization mechanisms → adds more complexity

TRANSPARENCY

Fault transparency

- **Hides errors and failures** (software/hardware) from the user
 - i.e. Automatic failover when resource is down
- Relates to access/replication transparency
 - Automatic relocation to replica
- May be not possible under certain assumptions
 - How to differentiate between timeout and dead resource?
 - Network access may experience problems at any time

TRANSPARENCY

- Is **total transparency** always necessary/preferable? **No**
- For instance, location transparency may be impossible if resource are geographically very distant.
- Tradeoff Transparency/Performance.
 - **Masking failures might result in slowness.**
- **Context-awareness** depending on location.
- Consider tradeoffs/consequences of transparency when designing a distributed system.

TRANSPARENCY

- Use of the word 'transparency' might be misleading.
- From the system's global point of view:
 - System as a whole appears as a *black box* to the user → 'opacity'
- The term 'transparency' is normally used from the *user's* perspective to talk about *parts* of the system
 - The user interacts with the system in some way.
 - A *change* is introduced in the system's behaviour (i.e. a call that was previously local now is remote).
 - The change is *transparent* to the user because all the user sees is exactly the same as before → the change is 'invisible'.
- 'Transparency' in distributed systems refers to distribution complexity as *part* of the system




OPENNESS

- **Expose** services according to consistent **syntax and semantics**.
- Interface definitions between components.
 - Semantics usually given in natural language.
- Specification **completeness and neutrality**.
- Completeness: everything needed is specified.
- Implementation neutrality.
 - Compatibility of implementations of different parties.
- **Interoperability and portability**.
- **Extensibility** / Replacing components.
- Example: RPC interface definition, Network File System (NFS)




SCALABILITY

- If we add more users, can we easily add more resources to the system?
- Problems:
 - **Centralized** services.
 - Centralized data.
 - Centralized algorithms (i.e. routing)
- Decentralized algorithms
 - No machine has complete information.
 - Decisions made from local information.
 - Failure tolerant.
 - No assumption about global time.




SCALABILITY

- How to achieve geographical scalability?
- From LAN to WAN
 - Synchronous → Asynchronous
 - Service location via broadcast not possible
- Communication reliability can suffer.
- More network latency!



SCALABILITY

- Administrative scalability.
- Necessity to expand to another (administrative) domain.
- Adapt to new policies
 - Resource usage.
 - Payment .
 - Security
 - Protection against new domain.
 - Protection new domain again distributed system.



HOW TO SCALE?

I. Hide latencies

a. From **synchronous to asynchronous** communication

```
// Synchronous block

try {
    int result = SendRequest(params);

    // wait until request is completed
}
catch (Exception e)
{
    // Timeout? Abort
}
```

```
// Asynchronous block

Int result = new Thread {doTask(params)};

DoOtherStuff();

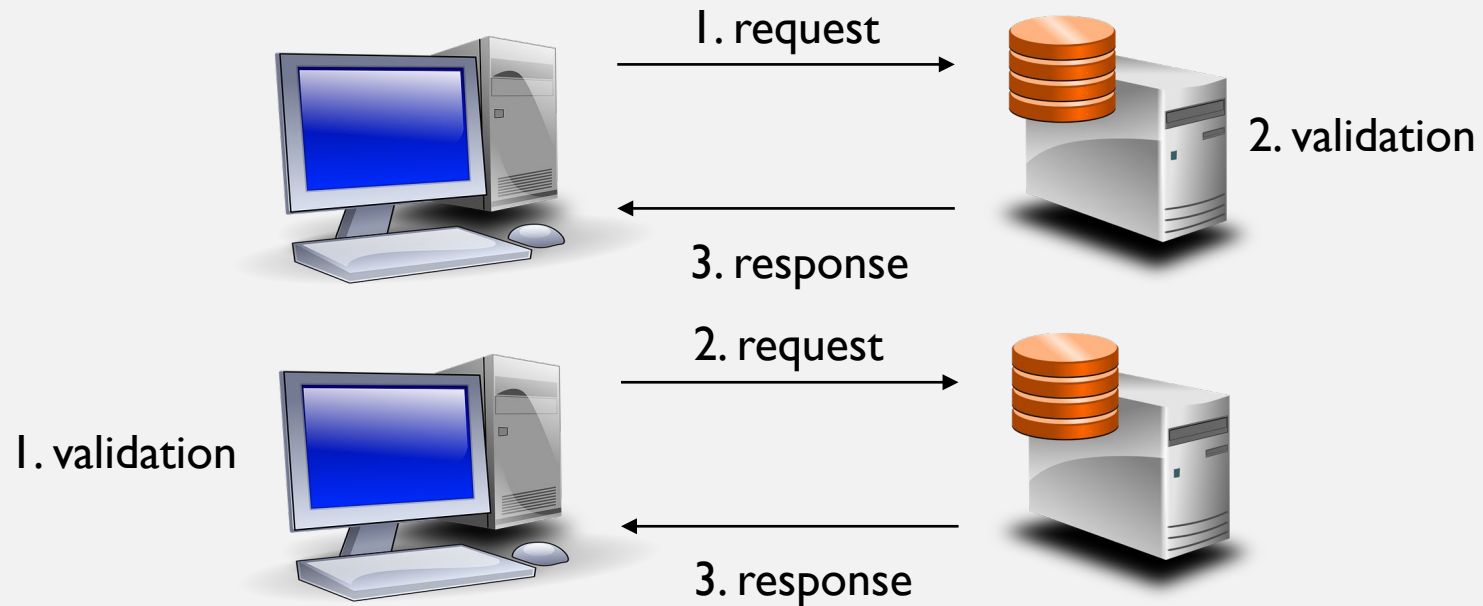
// in another thread

int doTask(Params params) {
    int result = SendRequest(params);
    return result;
}
```

HOW TO SCALE?

I. Hide latencies


b. Move functionality to the client



HOW TO SCALE?

2. Distribution


- Distribute components over the system
- e.g. DNS
- Naming resolution is distributed among many servers
- No centralized name resolution on a single server



HOW TO SCALE?

3. Replication


- Store replicas of a resource nearby
- How to modify a copy?
- **Temporary inconsistency** (i.e. when caching a document)
- If no inconsistencies allowed → **global synchronization**
 - **Very complex**
 - May make scaling impossible



SOME PITFALLS

- ***What could (possibly) go wrong?***

1. The network is reliable.
2. The network is secure.
3. The network is homogeneous.
4. The topology does not change.
5. There is no latency.
6. Bandwidth is not a problem.
7. Transport cost is negligible.
8. There is only one administrator.




TYPES OF DISTRIBUTED SYSTEMS

1. Distributed **computing** systems
 - a. Cluster computing
 - b. Grid computing
2. Distributed **information** systems
 - a. Transaction processing systems
 - b. Enterprise application integration
3. Distributed **pervasive** systems
 - a. Ubiquitous computing
 - b. Sensor networks

DISTRIBUTED COMPUTING SYSTEMS

Cluster computing

- Cheaper to build **supercomputing capabilities** by combining off-the-shelf hardware instead of a single workstation.
- **Homogeneous hardware.**
- **High speed network.**
- Run a program on **parallel** in several machines.
- Example: Apache Hadoop.



rack = block of trays with computers

DISTRIBUTED COMPUTING SYSTEMS

Grid computing

- Heterogeneous environment
 - PCs, workstations, supercomputers, clusters, etc.
 - Loose coupling
- Aims to solve large scale computing problems.
- No assumptions on hardware, operating system, network, etc.
- Brings together resources from different organizations.
 - Large pool of computing resources.

same prog on each computer that executes in idle time
mix computing power systems
not same network

DISTRIBUTED INFORMATION SYSTEMS

- **Integration** of existing applications/services
 - E.g. enterprise environment
- Main problem: How to propagate changes from the client to the applications, or between the applications themselves?
- Possible solutions
 - Use of **distributed transactions**.
 - Use of **communication middleware**.
 - Remote Procedure Calls (RPC)
 - Remote Method Invocation (RMI)
 - Message-Oriented Middleware (MOM)

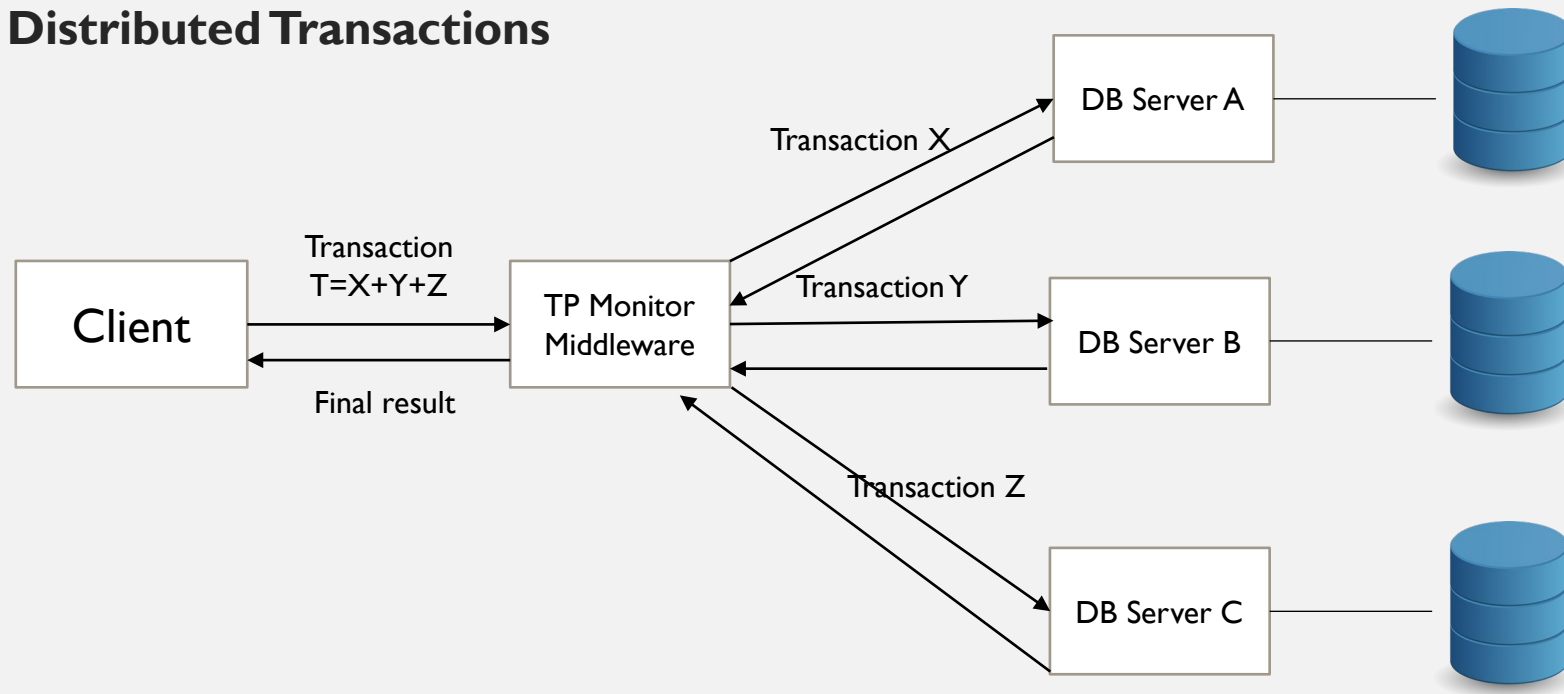
DISTRIBUTED INFORMATION SYSTEMS

Distributed transactions

- Single transactions according to **ACID**: Atomic, Consistent, Isolated, Durable |
- Atomic: either completely, or not at all.
- Consistent: invariants of the system are not changed.
- Isolated: transactions are run in (some) sequential order.
- Durable: changes are permanent and cannot be undone.

DISTRIBUTED INFORMATION SYSTEMS

Distributed Transactions



DISTRIBUTED INFORMATION SYSTEMS


Distributed transactions

- TP middleware monitors the transaction
- Transaction is split into nested transactions
- In order to make Y visible to X, X must be executed completely
- If Y fails, X must be undone by the TP middleware
- Example: IBM Customer Information Control System (CICS)
 - Airline systems, ATM, insurance, etc.

DISTRIBUTED INFORMATION SYSTEMS

Enterprise Application Integration

- Direct communication between applications by means of communication middleware



DISTRIBUTED INFORMATION SYSTEMS

Enterprise Application Integration

- Different types of communication middleware
 - RPC
 - RMI
 - Message-based
 - Web-services


DISTRIBUTED PERVASIVE SYSTEMS

- Mobile and embedded computing resources.
- Small, battery-powered, wireless.
- Part of our surroundings.
- Must react to a changing environment.
 - e.g. Wireless connection available.
 - Service discovery.
- Include means of exchanging and sharing information.
- Smart Systems.



UBIQUITOUS SYSTEMS

- System is *pervasive* and *continuously present*.
- User *interacts continuously* with the system.
- User *might not* be *aware* of the *interaction*.
 - e.g. providing input for sensor networks.



Source: Khan Academy

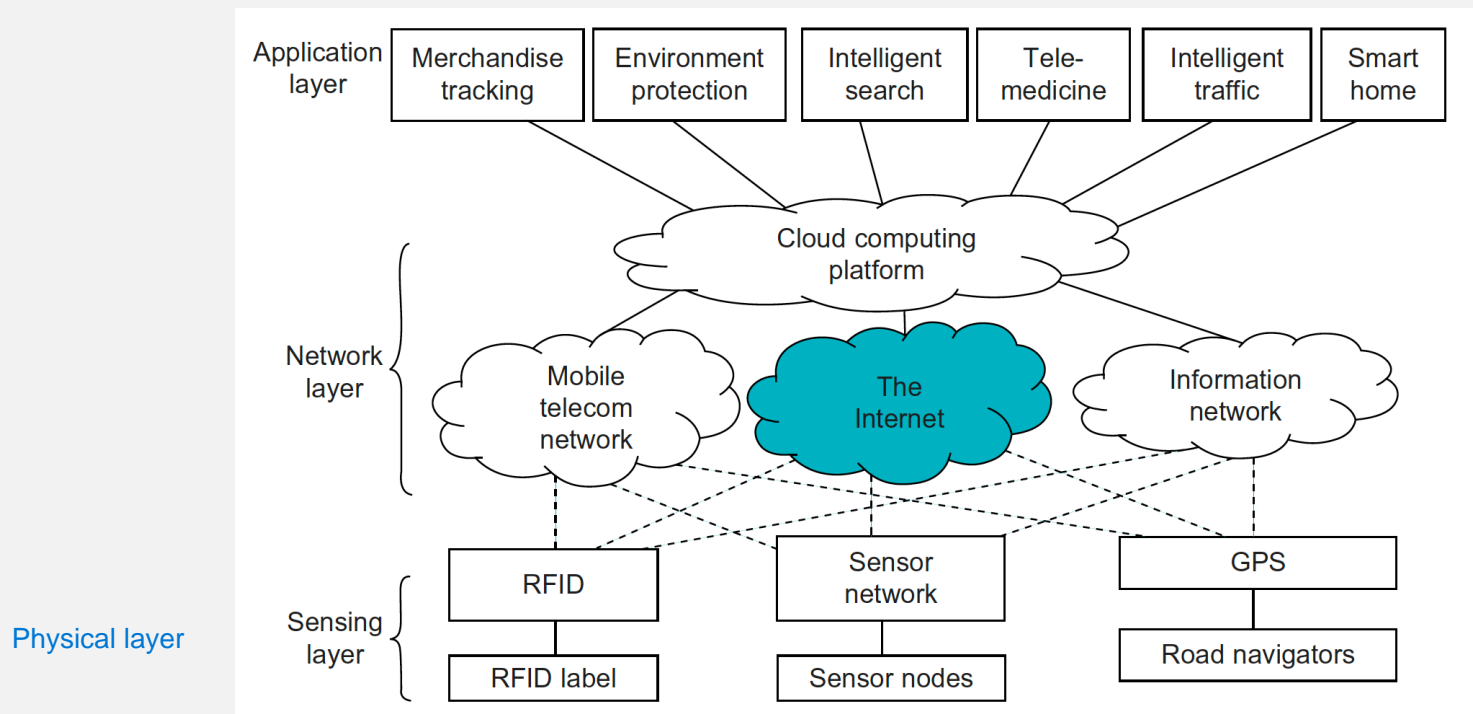
UBIQUITOUS SYSTEMS

Core requirements

- **Distribution**: devices are accessible in a transparent manner.
- **Interaction**: minimally intrusive.
- **Context awareness**: where, who, when and what.
 - High abstraction level so that applications can use those data.
- **Autonomy**: minimum administrative effort.
 - Joining/leaving the network, configuration, discovery, updates.
- **Intelligence**: React to changing situations by advanced algorithms.

UBIQUITOUS SYSTEMS


Internet of Things (IoT)



UBIQUITOUS SYSTEMS

Internet of Things (IoT)


- Many devices connected via radio frequency identification (RFID)
 - Sensor networks
 - Home systems
 - Road navigators
- RFID
 - Devices connected by radio waves
 - RFID tags can be read by RFID readers
 - Active: tag can send and receive information by itself (battery powered)
 - Semi-active: relies on reader for broadcasting (battery powered)
 - Passive: relies on reader completely (no battery)



DISTRIBUTED PERVASIVE SYSTEMS


Internet of Things (IoT)

- **Smart Buildings**
 - Residential, commercial, industrial, government settings
 - Smart homes, hospitals, office towers, shopping malls, etc.
 - **Monitoring** an regulation of heating, air conditioning, lightning, changes in the environment
 - **Security**, fire alarms, elevators



SENSOR NETWORKS

- **Monitor** and **record** conditions at different locations in space
- Wind, temperature, humidity, traffic, body-functions, etc.
- Lightweight/**portable sensor nodes**
 - **Transducer**: generates signals based on sensed information
 - **Microcomputer**
 - **Transceiver**: receives commands and exchanges data from/to central computer
 - **Power source**



SENSOR NETWORKS

- Network self-organization: unfeasible to configure each sensor node manually.
- Nodes may fail.
- New nodes may join the network.
- Collaborative signal processing.
 - Fuse data from multiple nodes.
 - Transmission of data and control messages.
 - Local sink nodes collect data and creates summary messages.