

NAMING

Distributed Systems

4. Sem BSc Informatics

IMC FH Krems

LECTURE OVERVIEW

- Introduction and motivation
- Flat naming
- Structured naming
- Attribute-based naming

NAMING

Why are names important in distributed systems?




NAMING

Types of names

1. Address: point of access where the entity can be found.
 - Example: network address of a machine.
2. Identifier: unique name used for an entity.
 - An identifier refers to at most one entity.
 - Each entity is referred to by at most one identifier.
 - An identifier always refers to the same entity (will not be reused for another entity).
 - Example: key in a distributed hash table (DHT).
3. Human-friendly names: character strings readable by humans.
 - Example: domain names, file names.

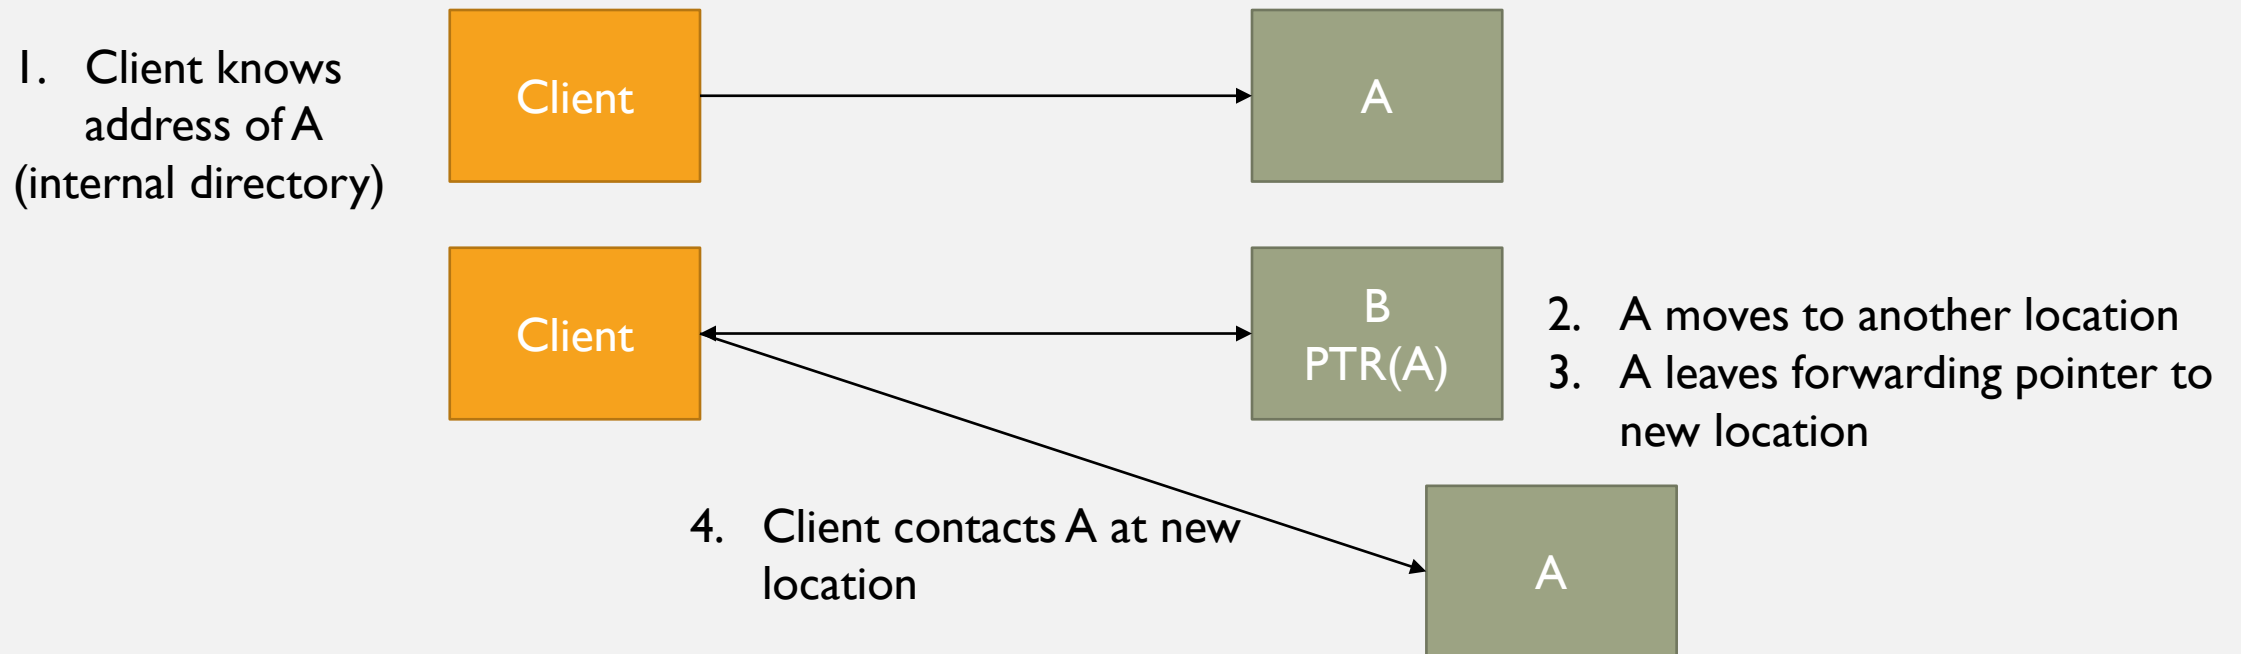
FLAT NAMING

- Resolves the mapping identifier → address.
- First approach: broadcasting
 - Identifier is broadcast to each machine in a network.
 - Machines with an entry point for that entity send reply.
 - Example: **Address Resolution Protocol (ARP)**.
 - IP address is sent to all machines of a network via broadcast message.
 - Machine with that IP address replies with its network address.
 - Only efficient on moderate-size networks (e.g. LAN).



FLAT NAMING

- Second approach: forwarding pointers.



FLAT NAMING

- Drawbacks of forwarding pointers
 1. Forwarding chains can become very long.
 2. Each node has to be maintained as long as necessary.
 3. Broken links.
- Chains have to be kept short and up to date → maintaining effort.


FLAT NAMING

- Next approach: **Distributed Hash Tables (DHT)**.
- (key, value) pairs are distributed among the nodes of the network.
- Value can be an address, or the entity itself (e.g. a document).
- Specially important in P2P networks.
 - Examples: BitTorrent, Freenet.
- Nodes store a part of the lookup table with a subset of entries.
- Nodes that join/leave the network are updated locally.
 - No central coordination.
- Extremely scalable approach.

FLAT NAMING

Distributed Hash Tables

- Example: **Chord**.
- Identifier space in a m -bit space (e.g. 160 bits)
 - SHA-1 hash of a name is calculated as key.
 - Potentially 2^m nodes forming a circle.
- Both **nodes** and **entities** have identifiers in this space.



FLAT NAMING

Chord

- Entity with id k is assigned the node with the smallest id such that $id \geq k$. This node is called the *successor* of entity k , $\text{succ}(k)$.
- Our problem is to find the address of $\text{succ}(k)$, given k .
- Each node stores a *finger table* (FT) with entries pointing to other nodes.
 - Entry j points to the first node succeeding the current node by at least 2^{j-1}
 - Example: $\text{FT}[1] \rightarrow$ points to first node q succeeding the current node by at least 1
 - $\text{FT}[2] \rightarrow$ points to node q' succeeding the current node by at least 2
 - $\text{FT}[5] \rightarrow$ points to node q'' succeeding the current node by at least 16
- Node p forwards request to node q in its finger table, so that it is the node with the smallest id that could store k .


FLAT NAMING

Chord: Example

- $m = 4$, currently 3 nodes

$\text{keys}(5) = \{1, 2, 3, 4, 5\}$

| | |
|---|---|
| 1 | 9 |
| 2 | 9 |
| 3 | 9 |
| 4 | 0 |



| | |
|---|---|
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 5 |

$\text{keys}(9) = \{6, 7, 8, 9\}$

| | |
|---|---|
| 1 | 5 |
| 2 | 5 |
| 3 | 5 |
| 4 | 9 |

$\text{keys}(0) = \{10, 11, 12, 13, 14, 15, 0\}$

FLAT NAMING

Chord: Example

- Resolve key = 6

3. Lookup $k = 6$

4. Forward to node 9

$\text{keys}(5) = \{1, 2, 3, 4, 5\}$

| | |
|---|---|
| 1 | 9 |
| 2 | 9 |
| 3 | 9 |
| 4 | 0 |

5

0

9

| | |
|---|---|
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 5 |

$\text{keys}(9) = \{6, 7, 8, 9\}$

5. Lookup $k = 6$

6. Found in node 9

1. Lookup $k = 6$

2. Forward to node 5

| | |
|---|---|
| 1 | 5 |
| 2 | 5 |
| 3 | 5 |
| 4 | 9 |

$\text{keys}(0) = \{10, 11, 12, 13, 14, 15, 0\}$

FLAT NAMING

Chord


- Join the network: new node p looks up $\text{succ}(p+1)$ and inserts itself

$\text{keys}(5) = \{1, 2, 3, 4, 5\}$

| | |
|---|---|
| 1 | 5 |
| 2 | 5 |
| 3 | 9 |
| 4 | 0 |

1. $\text{succ}(4) = 5$
2. Node 3 inserts itself
3. $\text{keys}(3) = \{1, 2, 3\}$

| | |
|---|---|
| 1 | 9 |
| 2 | 9 |
| 3 | 9 |
| 4 | 0 |



| | |
|---|---|
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 3 |

$\text{keys}(9) = \{6, 7, 8, 9\}$

3. Rest of nodes issue $\text{succ}(k)$ requests regularly, updating their FT accordingly

| | |
|---|---|
| 1 | 3 |
| 2 | 3 |
| 3 | 5 |
| 4 | 9 |

$\text{keys}(0) = \{10, 11, 12, 13, 14, 15, 0\}$

FLAT NAMING

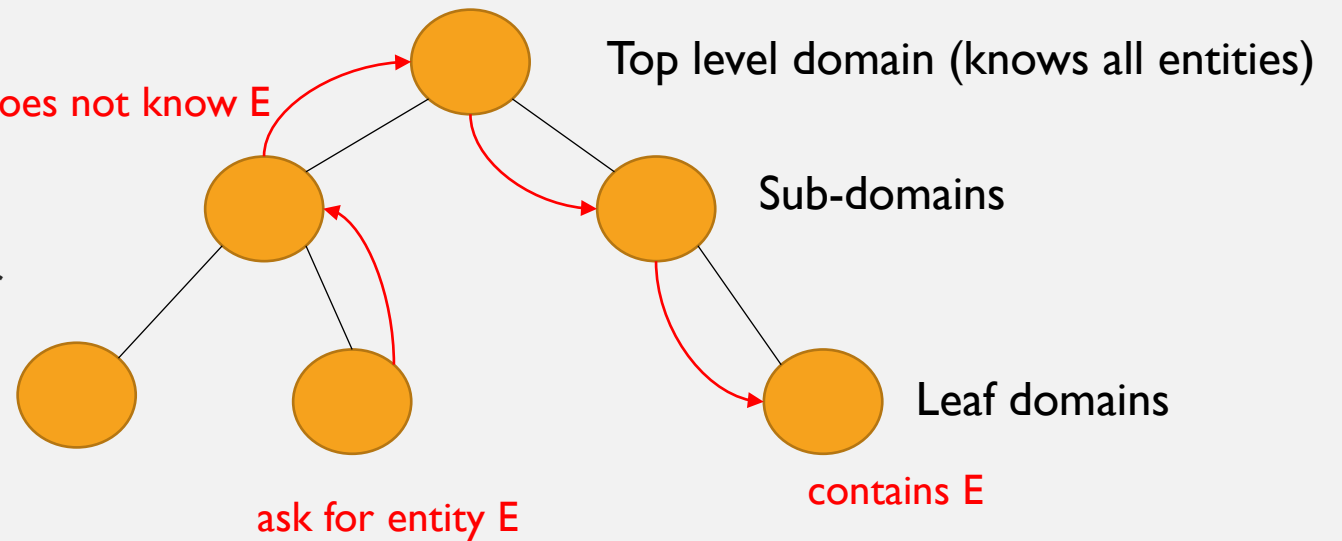
- Last approach: hierarchical naming

- Based on *domains*.

- Each domain contains entities. **does not know E**

- When looking up an entity

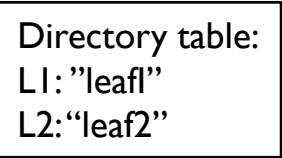
- a. Entity is not there → ask higher level domain.
- b. Entity is there → address of entity itself or pointer to sub-domain.



STRUCTURED NAMING

- Flat names useful for machines, but normally not meaningful for humans.
- How to generate unique names that are human readable?
- Solution: use name spaces.
- Name spaces can be visualized as name graphs (labeled directed graphs) with:
 - Root node: node without ingoing edges.
 - Directory nodes: nodes with outgoing edges → stores directory tables.
 - Leaf nodes: nodes without outgoing edges → entities.
 - Value can be the address of the entity or the entity itself (e.g. a file).

STRUCTURED NAMING



STRUCTURED NAMING

- Name resolution
 - Begin with root node.
 - Lookup directory table to find the node N_1 the first label refers to.
 - Contact node N_1 and look up next label in directory table to find node N_2 .
 - Contact node N_2 and look up the next label in directory ...
 - ... until leaf node reached or entity non-existent

STRUCTURED NAMING

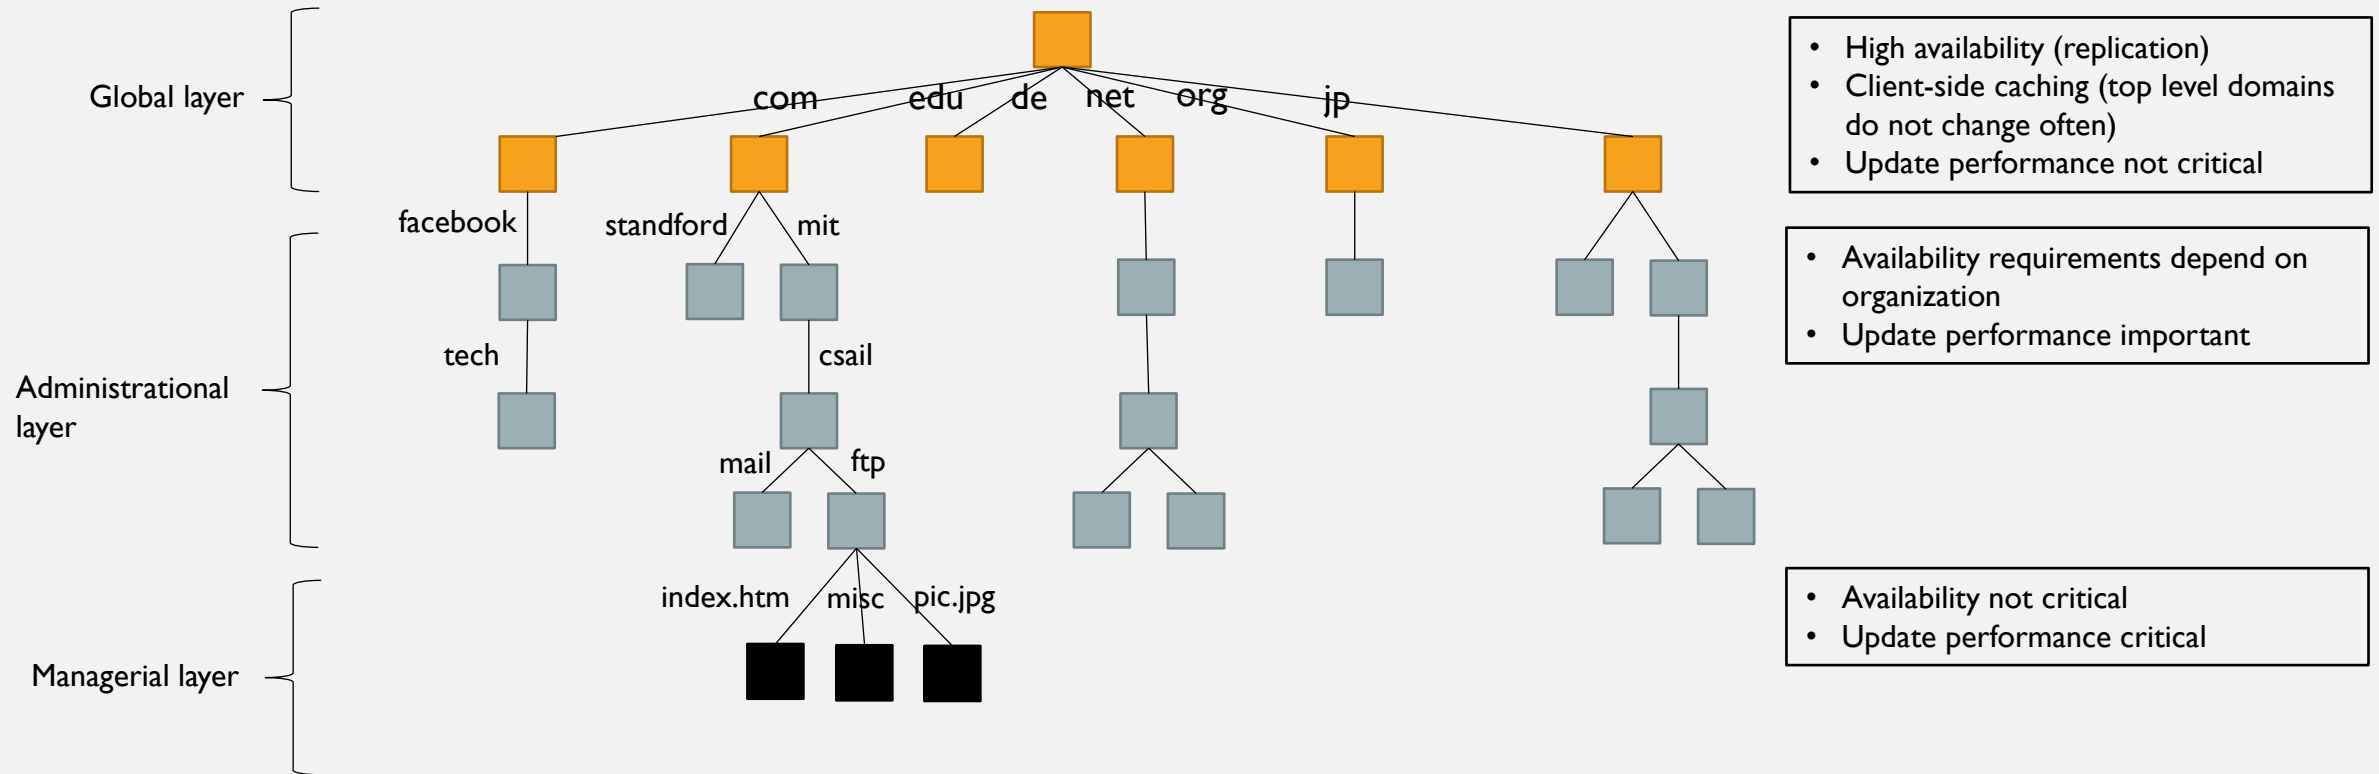
- Naming implemented in distributed systems by **name servers**.
- Number of name servers depending on geographical spread of the system.
- Distributed name spaces are organized into logical layers:
 - **Global layer:** root node and its direct children. Most stable layer (directory tables are rarely changed).
 - **Administrational layer:** directory nodes maintained by a single organization (i.e. a directory node per department inside an organization). Changes more than global layer.
 - **Managerial layer:** Nodes represent entities that regularly change (i.e. shared files or directories). Managed also by end-users.

STRUCTURED NAMING

- Example: **Domain Name System (DNS)**.
- Translation name (URL) → IP address.
- One of the backbone services of the Internet.
- Current 13 logical root name servers
 - Each one operates redundantly (high availability).
- Normally uses UDP on port 53.
- **Domain:** subtree of the DNS naming graph.
- **DNS Zone:** contiguous portion of the domain space which is administrated by a single entity (with a name server). A domain can be portioned into zones.

STRUCTURED NAMING

- Example: **Domain Name System (DNS)**



STRUCTURED NAMING

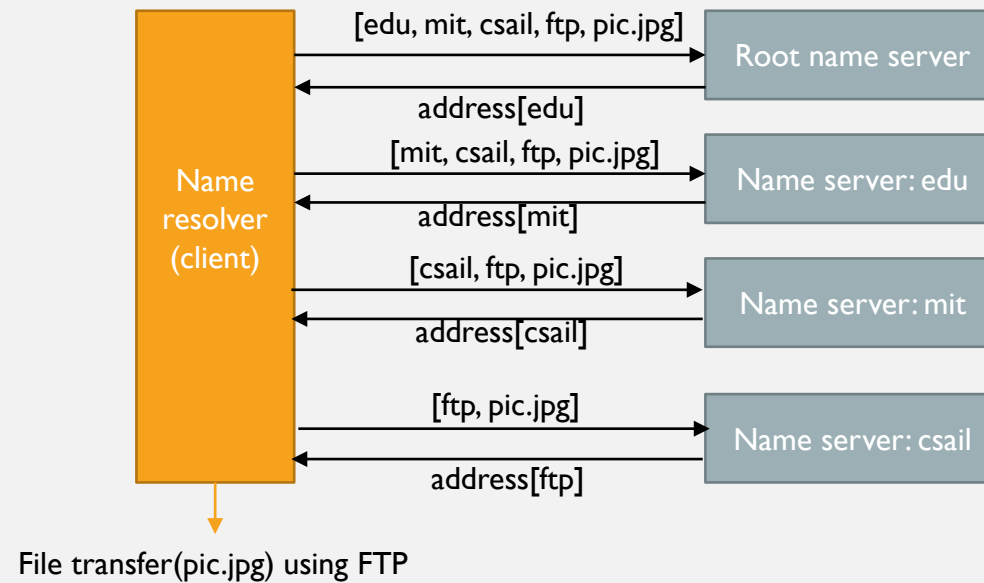
Domain Name Service (DNS)

- Each zone is implemented by a name server.
- A name server contains a zone file with records for all nodes in a zone.
- DNS record types: relate a domain name with certain data
 - SOA (Start of authority): metadata about the zone (administration).
 - A (address): the IP address of a particular host.
 - NS (name server): the IP address of a name server implementing the zone.
 - MX (mail exchange): the IP address of a mail server.
 - SRV (service): record for a specific service.

STRUCTURED NAMING

DNS name resolution: [edu, mit, csail, ftp, pic.jpg] with ftp (ftp://ftp.csail.mit.edu/pic.jpg)

I. Iterative name resolution



STRUCTURED NAMING

DNS name resolution: [edu, mit, csail, ftp, pic.jpg] with ftp (ftp://ftp.csail.mit.edu/pic.jpg)

2. Recursive name resolution

