

SYNCHRONIZATION

Distributed Systems

4. Sem BSc Informatics

IMC FH Krems

LECTURE OUTLINE

- Clock synchronization
- Logical Clocks
- Mutual exclusion
- Election algorithms
- Distributed event matching

SYNCHRONIZATION

- Remember definition of distributed systems from Lecture 1: “A *distributed system is a collection of independent computers that appears to its users as a single coherent system*”
- Question: **How** do distributed systems manage to appear to its users as a single coherent system?
- Processes in a distributed system **collaborate** to achieve this goal.
- **Communication** is a part answer to this collaboration problem.
- **Coordination** is the next answer to the question.
 - Process coordinate by means of synchronization mechanisms.

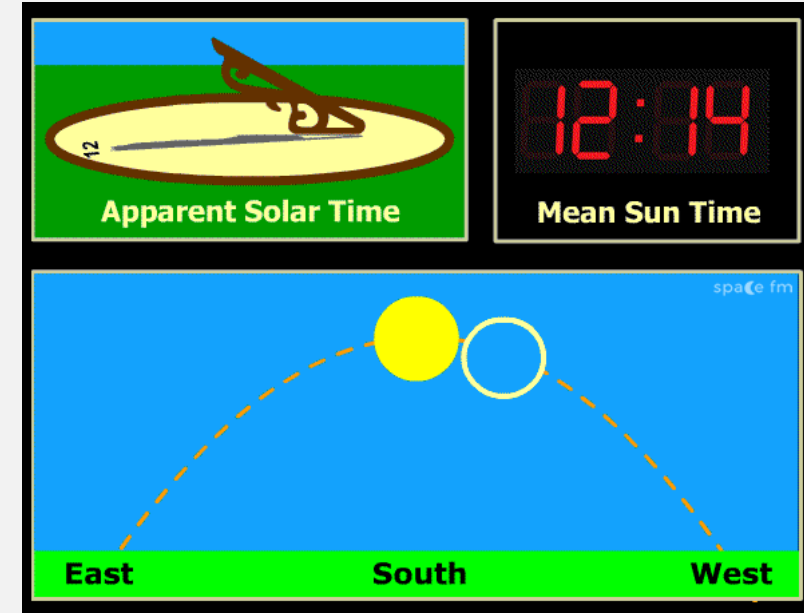
CLOCKS

- In a centralized system, time is not ambiguous.
- In a distributed system, there might be small time differences between distributed processes.
- We need agreement on *what happened before/after what*.
- Possible approaches:
 1. Synchronize all clocks in a distributed system.
 2. Use *logical clocks*.



CLOCK SYNCHRONIZATION

- Goal: synchronize **physical** clocks.
- Measurement base is **Universal Coordinated Time (UTC)**
 - Laboratories around the world measure time using Cs 133 clocks.
 - Bureau International de l'Heure (BIH - Paris) averages measurements to International Atomic Time (TAI).
 - From time to time, there is a clock skew between TAI and mean solar time.
 - When this happens, BIH announces a leap second.
 - $UTC = TAI + \text{leap seconds}$.
- UTC is broadcasted using shortwave radio stations.
- In combination with satellite services → accuracy ± 50 ns.



Source: space.fm


CLOCK SYNCHRONIZATION

Clock synchronization algorithms

- Goal: keep clocks in a distributed system synchronized within a given precision.
- Every hardware clock is subject to a **clock drift** by physical reasons.
 - Standard quartz clocks show drift rate of approx. 30 sec per year.
- **Network Time Protocol (NTP).**
 - Clients contact a time server with UTC receiver and accurate clock.
 - Reported time differs of server time by message delays.

CLOCK SYNCHRONIZATION

Network Time Protocol (NTP)



- Goal: estimate offset $\delta = T_2 - T_1$ (where we assume that $T_2 - T_1 \approx T_4 - T_3$)

$$\delta = \frac{(T_4 - T_1) - (T_3 - T_2)}{2}$$

- What happens if $\delta < 0$?

CLOCK SYNCHRONIZATION

Network Time Protocol (NTP)

- One of the oldest protocols in the Internet.
- UDP port 123.
- Can achieve accuracy in range 1-50 ms.
- Client (which can be another NTP server) synchronizes its clock.
 - Either $\delta > 0$: clock must be set forward.
 - Or $\delta < 0$: client's clock is set to a smaller frequency, until both clocks are in sync.
- Servers/clocks organized in *strata*.
 - Stratum 0: atomic clocks.
 - Stratum 1: those servers directly synchronized with an atomic clock (stratum 0).
 - In general, if server *A* synchronizes with a stratum *k* server, then *A* becomes a stratum *k* + 1 server.

LOGICAL CLOCKS

- For most applications, we don't need all clocks to be synchronized with true UTC time.
- It is enough for all machines in a (closed) system to agree about what events preceeded other, and vice-versa.
- Logical clocks, introduced by Lamport (1978) achieve this.
- Let a, b be two events in a distributed system.
- We define a relation Happens—before(a, b) denoted by $a \rightarrow b$ such that:
 1. If a and b originated in the same process, and a happened before b , then $a \rightarrow b$
 2. If a is the event of sending a message and b the event of receiving that message, then $a \rightarrow b$
- Events that happened before another event (according to this relation) are said to be *causally* related.


LOGICAL CLOCKS

- What we look for is a logical clock C such that all process agree that
$$\text{If } a \rightarrow b \text{ then } C(a) < C(b)$$
- The relation is transitive: if $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$
- If two events x and y happen in processes that do not exchange messages, then neither $x \rightarrow y$ nor $y \rightarrow x$ (x and y are concurrent).
 - C induces a partial order in the set of events.
- Corrections to a logical clock C can be made by adding values, never by subtracting.

this clocks are more like counters rather than clock in classical way

LOGICAL CLOCKS

- Each process P_i has a clock C_i running at a given frequency



- P_2 receives message m_1 from P_1 with timestamp $t_s(m_1) = 5$
- P_3 receives message m_2 from P_2 with timestamp $t_s(m_2) = 14$


LOGICAL CLOCKS

General algorithm for Lamport clocks:

- Process P_i has clock C_i which is incremented by 1 with a given frequency.
 - C_i can be seen as a general event counter.
- If process P_i sends a message m to process P_j , it sends as timestamp $t_s(m) = C_i$
- Upon receiving m , P_j updates its clock by $C_j \rightarrow \max(C_j, t_s(m))$.
- P_j increments its clock by 1 and forwards the message to the receiving application.
- Lamport clocks are implemented in the middleware layer.

LOGICAL CLOCKS


- Can we conclude from the counters that m_6 was **sent** before m_5 ?



VECTOR CLOCKS

- If $C(a) < C(b)$ we cannot always infer that a actually happened before b .
- For that, we need to account for the *causal histories* of the events that happened in the system.
- Main idea: extend Lamport clocks with the clocks of all the other processes in the system.
- Resulting construction: **vector clock**.
- This vector clock is sent along with each message.


VECTOR CLOCKS



VECTOR CLOCKS

- $VC_i = [C_1, C_2, \dots, C_i, \dots, C_n]$ is a vector of logical clocks, where $VC_i[i]$ is the current logical clock (= event counter) of process P_i .
- Upon an event occurring (i.e. before sending a message) $VC[i] \leftarrow VC[i] + 1$
- If P_i sends a message to P_j then P_i sends its vector clock along with the message.
- Upon receiving the message, P_j adjusts its own clock by $VC[k] \leftarrow \max(VC[k], t_s(m))$ for all processes k and then increments it .
- Comparing timestamps:
 - $a \rightarrow b$ if and only if $VC_a[k] \leq VC_b[k]$ for all k and there is at least one j such that $VC_a[j] < VC_b[j]$.

VECTOR CLOCKS



Are sending of m_4 and m_5 causally related?

LECTURE OUTLINE

- Clock synchronization
- Logical Clocks
- **Mutual exclusion**
- Election algorithms
- Distributed event matching

MUTUAL EXCLUSION

- How to ensure that processes that access the same resource simultaneously leave it in a consistent state?
- **Mutual exclusion:** ensure that only one process can access a given resource at a time.
- We need to prevent two common problems:
 1. **Starvation:** A process that needs to access a resource waits forever to get access granted.
 2. **Deadlock:** A situation where a number of processes wait for another process to finish using a resource, but that process is waiting as well.


MUTUAL EXCLUSION

I. Permission-based approach

- Based on logical clocks.
- Each process sends a request for a given resource as (process name, resource name, logical clock) to *all* other processes.
- When a process receives the request:
 - If it is not interested in the resource → do nothing.
 - If it is currently accessing the resource → queue the request and send a response.
 - If it is currently **not** accessing the resource, but wants to → compare its clock with the timestamp in the requests.
 - The lowest value wins and the request who “lost” gets queued.

MUTUAL EXCLUSION

I. Permission-based approach

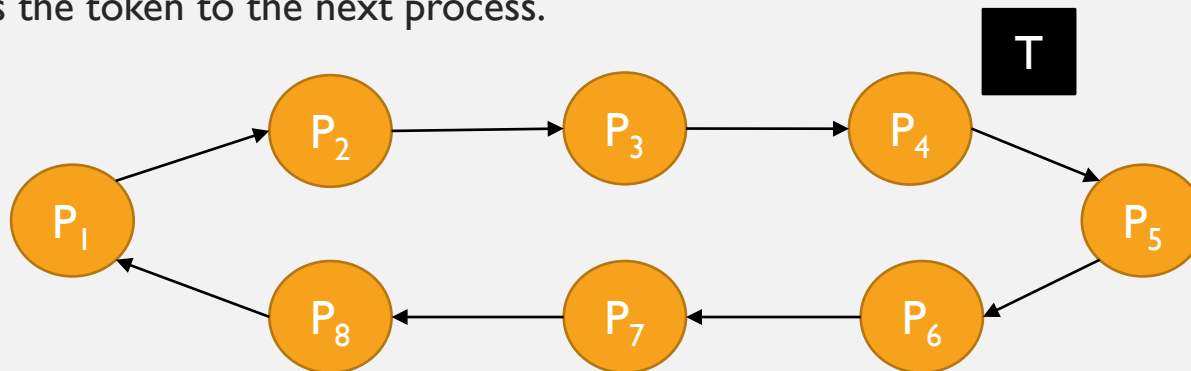


1. P_1 and P_3 are interested in accessing resource R
2. P_2 is not interested, so it ignores the request
3. P_4 is using the resource now
4. P_4 queues P_3 next
5. P_3 compares its clock to that of P_1 and queues it
6. When P_4 is done, it will handle the request to P_3
7. When P_3 is done, it will handle the request to P_1

MUTUAL EXCLUSION

2. Token-based approach

- Processes form an overlay network with the form of a ring.
- A token is generated and passed along the ring. The token is associated with a given resource.
- When a process receives the token, it might
 - Use the associated resource.
 - Pass the token to the next process.



LECTURE OUTLINE

- Clock synchronization
- Logical Clocks
- Mutual exclusion
- **Election algorithms**
- Distributed event matching

ELECTION ALGORITHMS

- In some situations, one process has to act with a special role, for instance, as a coordinator
 - For instance, only one process has rights to change some data.
 - Coordinator role for replicating data.
 - If coordinator fails, a new coordinator must be “appointed”.
- This process can be elected by an administrator, or automatically by means of an **election algorithm**.
- Main goal: after the algorithm, all processes should agree on which process has the coordinator role.

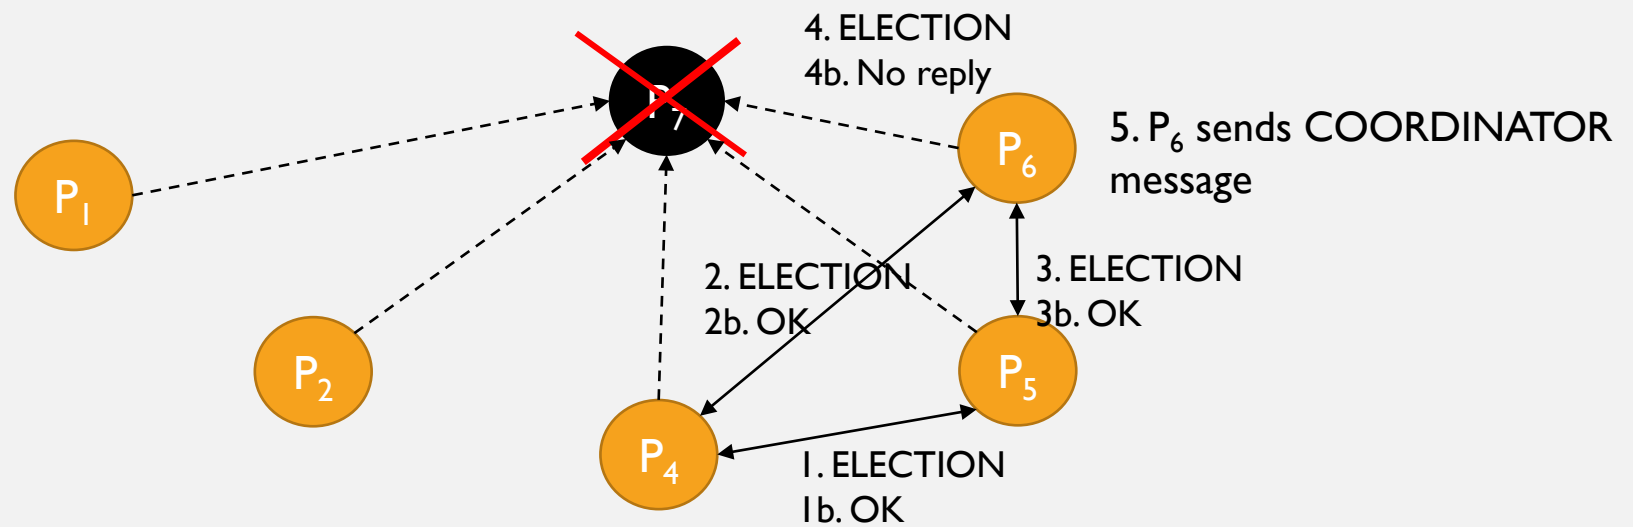
ELECTION ALGORITHMS

Bully algorithm

- A group of processes $\{P_1, P_2, \dots, P_N\}$ needs to elect a coordinator. We assume that processes are ordered in ascending order by their IDs $\text{id}(P_k) = k$
- A process k announces that it will hold an election by means of a ELECTION message and sends it to all processes above it P_{k+1}, \dots, P_N
- If some process $j > k$ answers, then it takes over and P_k waits for the result.
- If no process replies, then P_k has won. It will be the new coordinator and announces it by means of a COORDINATOR message to all other processes.

ELECTION ALGORITHMS

Bully algorithm



What should happen if P_7 recovers?

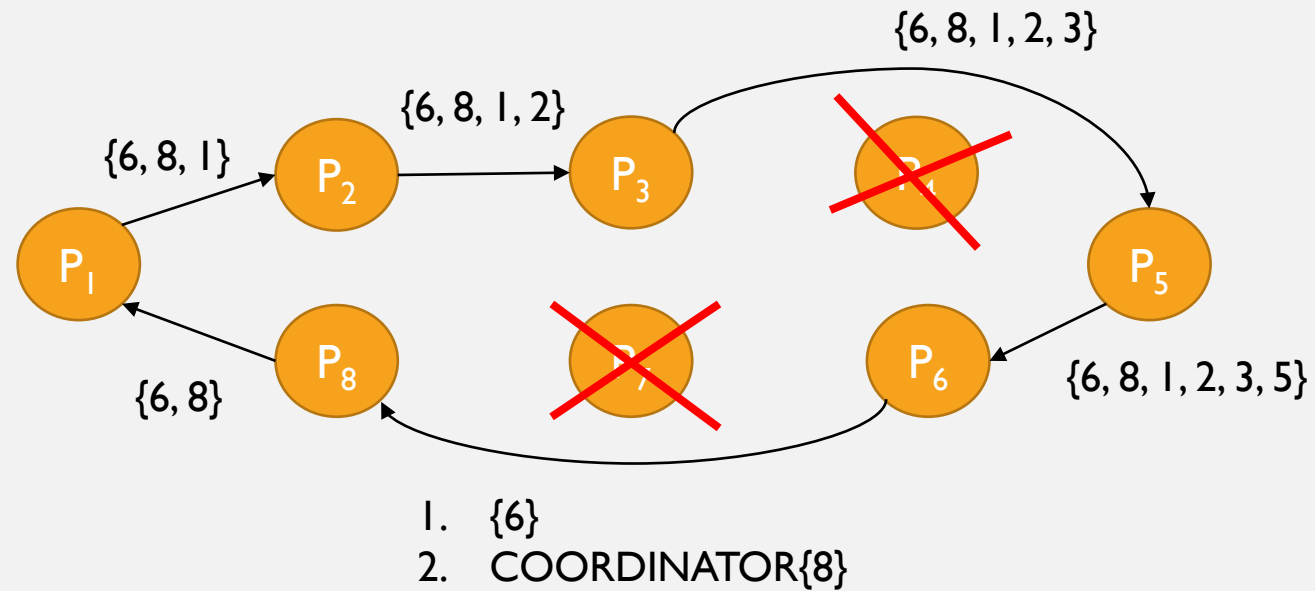
ELECTION ALGORITHMS

Ring algorithm

- Processes are arranged in a ring-like logical network.
- When coordinator fails, a process sends an ELECTION message, containing a list of process IDs.
 - The process includes itself in the list.
 - It sends the list to the next running process.
- Each process in the ring adds itself to the list.
- As soon as the message gets back, it changes into an COORDINATOR message to announce the new coordinator.
 - The new coordinator is the process with the highest ID in the list.

ELECTION ALGORITHMS

Ring algorithm



LECTURE OUTLINE

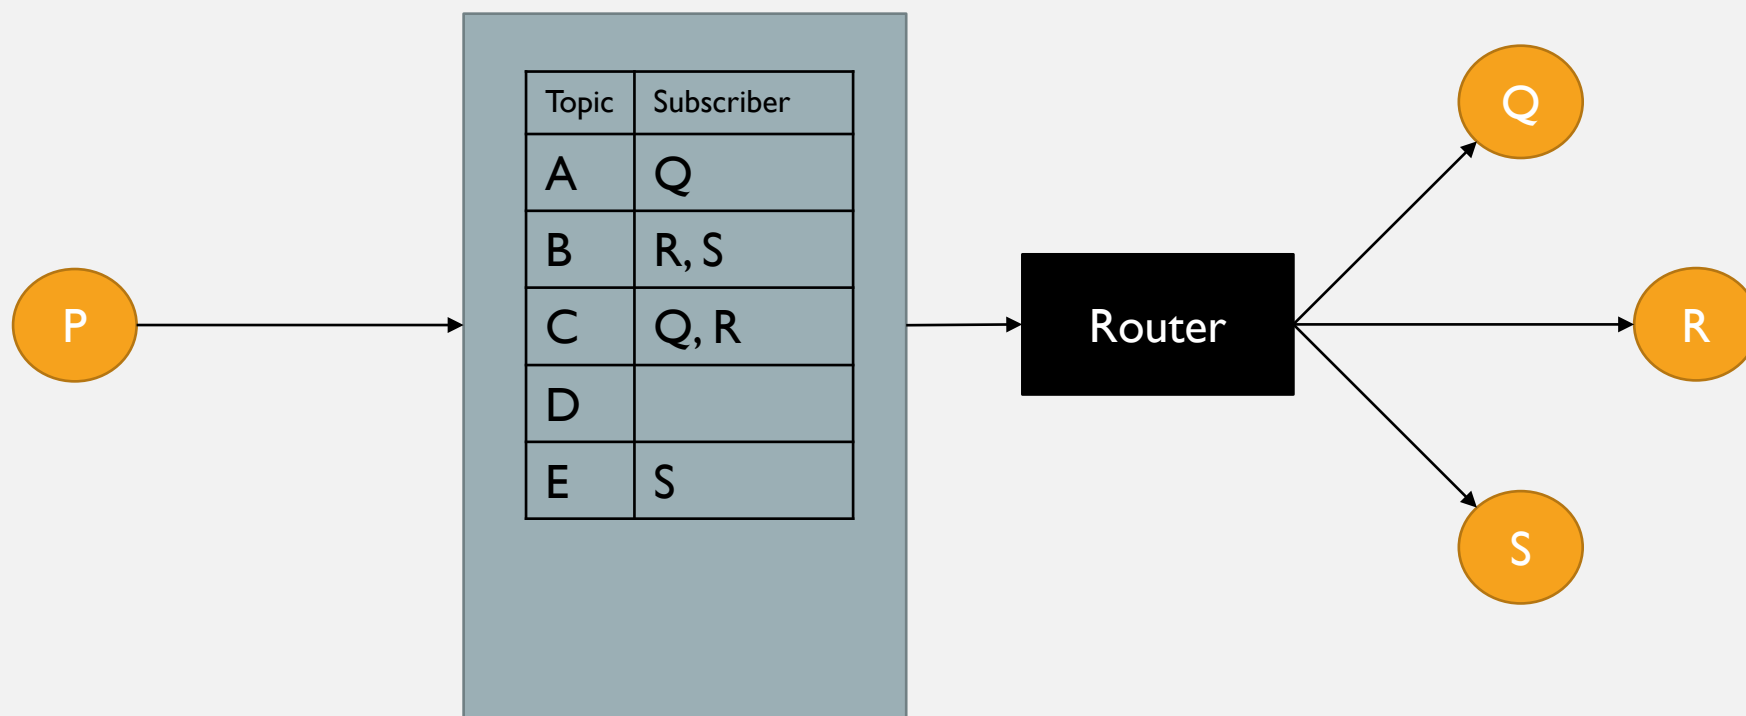
- Clock synchronization
- Logical Clocks
- Mutual exclusion
- Election algorithms
- **Distributed event matching**

EVENT MATCHING

- Remember **publish/subscribe** systems: a subscriber process P subscribes to a topic to get an event notification as soon as an event happens that is sent to that topic.
- Two main (coordination) issues:
 1. How to match subscriptions to events.
 2. How to notify subscribers in case of a match.
- Assume there is a function $\text{match}(S, N)$ that returns true if a subscription matches a notification, or false otherwise.

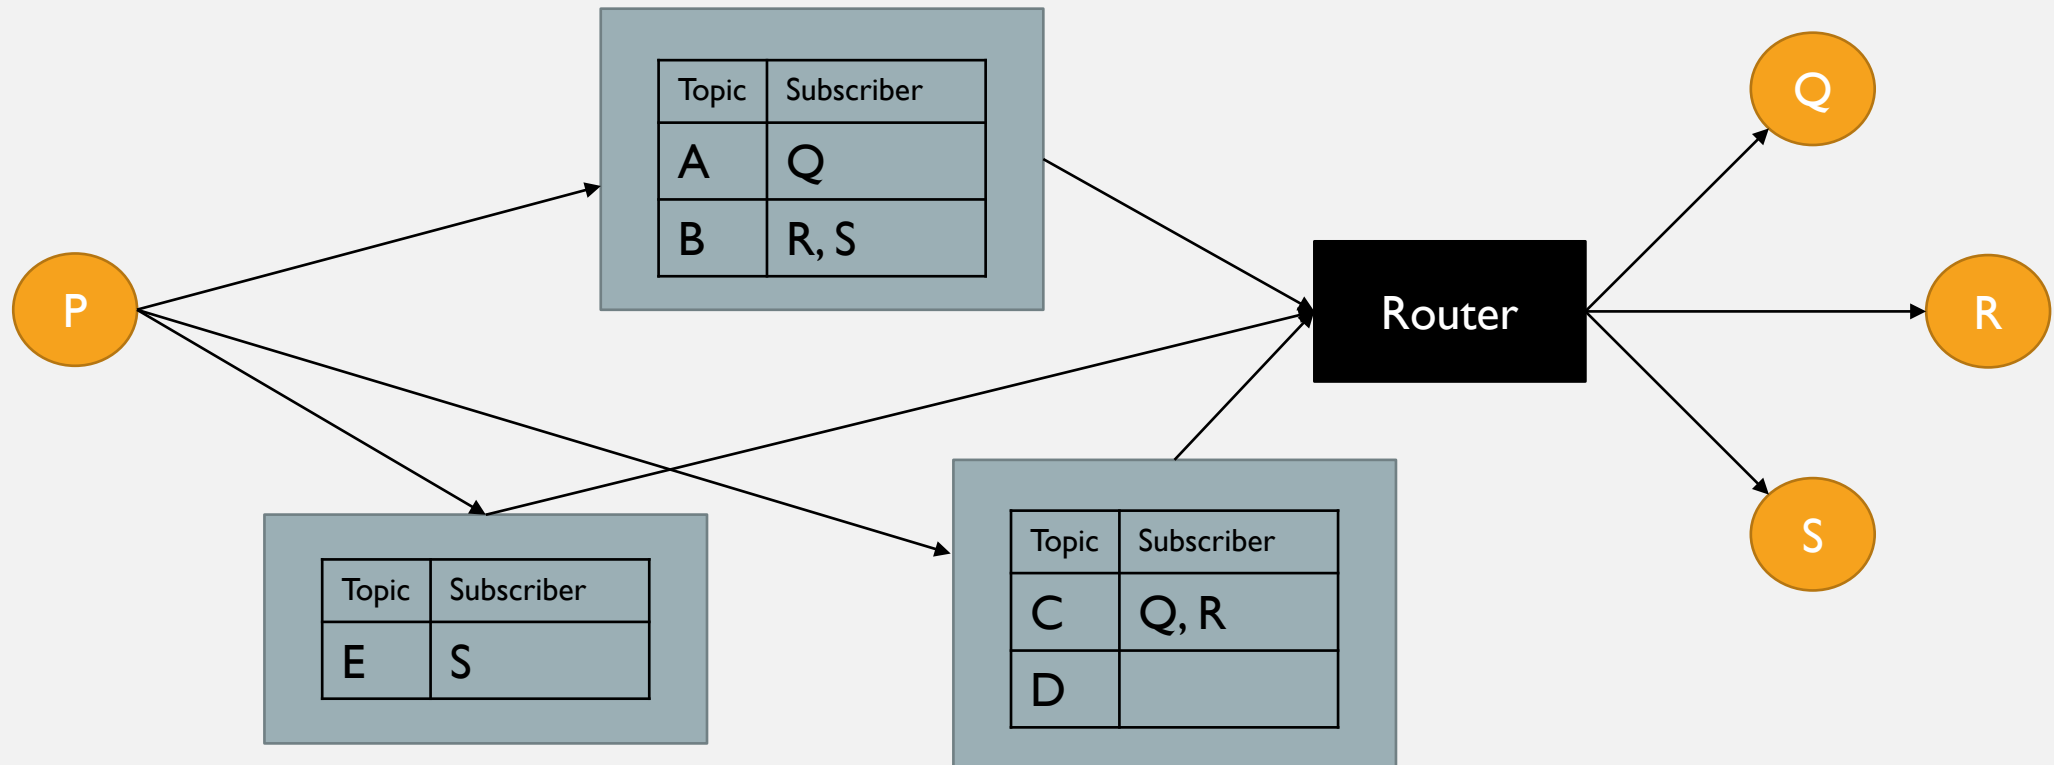
EVENT MATCHING

- Centralized implementation



EVENT MATCHING

- Distributed implementation: we divide the work among several servers



EVENT MATCHING

- Distributed implementation: we divide the work among several servers
- Subscriptions distributed among a set of servers (**brokers**)
 - Which server can handle which notification → can be solved by DHT.
 - Routing servers (**routers**) forward notification to subscribers.
- How to forward notifications to subscribers
 - Flooding (i.e. multicast).
 - Selective-notification routing.

EVENT MATCHING

- Distributed implementation: we divide the work among several servers
- Selective notification routing:
 - Server routes a notification based on (part of) its contents.
 - Depending on this information, some routes are discarded.
 - Each broker broadcasts its subscription over the (overlay) network.
 - Routers compose **routing filters**.

EVENT MATCHING

- Distributed implementation: selective notification routing

